

Recovering OpenSSL ECDSA Nonces Using the Flush+Reload Cache Side-channel Attack

No Author Given

No Institute Given

Abstract. In this paper we illustrate a vulnerability introduced to elliptic curve cryptographic protocols when implemented using a function of the OpenSSL cryptographic library. For a given implementation using an elliptic curve E over a binary field with a point $G \in E$, our attack can recover the majority of the bits of a scalar k when kG is computed using the OpenSSL implementation of the Montgomery ladder. For the Elliptic Curve Digital Signature Algorithm (ECDSA) the scalar k is intended to remain secret. Our attack recovers the scalar k and thus the secret key of the signer and would therefore allow unlimited forgeries. This is possible from snooping on only one signing process and requires computation of less than one second on a quad core desktop when the scalar k is around 571 bits.

1 Introduction

Elliptic curve cryptography (ECC) [11, 13] is a collection of public key cryptographic methods that rely on the computational intractability of the Elliptic Curve Discrete Logarithm Problem (ECDLP). In a nutshell, given an elliptic curve over a finite field and two points on the curve G and H , the ECDLP is to find the scalar k such that $H = kG$.

ECC offers a higher encryption strength per key-bit than related methods with security reliant on the hardness of computing discrete logarithms in finite field or factoring the product of large primes. Consequently, ECC uses significantly shorter keys and offers faster operations than other methods, contributing to its rising popularity.

The Elliptic Curve Digital Signature Algorithm (ECDSA) [1, 9, 15] is a standard digital signature algorithm based on elliptic curves. One core operation of the ECDSA algorithm, as in many ECC protocols, the scalar multiplication of a point on the elliptic curve by a pseudo-randomly generated secret nonce. The confidentiality of the nonce is paramount for the security of the algorithm. Past research indicates that partial exposure of nonce bits can be exploited for efficient attacks on the secret key [3, 17].

OpenSSL [19] is a cryptographic software package that implements ECDSA. When using elliptic curves over a binary field \mathbb{F}_{2^m} , OpenSSL uses the Montgomery ladder [10, 14] algorithm for multiplying the point by the nonce. One of the advantages of the Montgomery ladder is that it has regular behaviour, performing the same sequence of operations for each nonce bit, irrespective of the

value of the bit. This regular behaviour makes it more resilient to side-channel attacks [10, 18].

While the operations performed by the algorithm are regular, their targets depend on the value of the bits of the nonce. To apply the operations to the respective targets, the OpenSSL implementation uses a conditional branch based on the value of the bit. By tracing this branch an attacker can recover the values of the nonce bits and, consequently, break the cryptosystem.

In this paper we present our use of the FLUSH+RELOAD cache side-channel attack [22] to trace the branch in the OpenSSL implementation. FLUSH+RELOAD relies on a security weakness in the IA-32 and X86-64 architectures that allows processes to monitor other processes' read and execute access to shared memory pages. Our attack program monitors access to both arms of the conditional branch and uses the information collected from these probes to reconstruct the nonce. This attack is a threat to the security of any cryptographic protocol using this scalar multiplication method, when the attacker has access to the target computer's memory. In this paper we illustrate the efficiency of the attack by analysing ECDSA and recovering the secret key using only one signature at very little computational cost (in both time and memory).

The paper also presents new information on the limitation of the FLUSH+RELOAD attack. We discuss spatial limitations, affecting the distance between multiple probes, and temporal limitations, affecting the probe resolution.

The results of this paper also support the findings of [21] that longer keys render a cryptographic algorithm more vulnerable to side-channel analysis.

The rest of this paper is organised as follows. The next section presents background information on elliptic curves, ECDSA, the Montgomery ladder and the FLUSH+RELOAD attack followed by a short discussion of related research. Section 3 describes our attack on the OpenSSL implementation of ECDSA. The results of the attack are analysed in Section 4. We discuss the implications of the attack and suggest techniques for mitigation in Section 5.

2 Preliminaries

In this section we set the scene for our attack.

2.1 ECDSA

The ElGamal Signature Scheme is the basis of the US 1994 NIST standard, Digital Signature Algorithm (DSA). The ECDSA is the adaptation of one step of the algorithm from the multiplicative group of a finite field to the group of points on an elliptic curve. The main benefit of using this group over field elements is smaller parameters for the same security level as mentioned above. ***

Parameters: An elliptic curve E defined over a finite field \mathbb{F}_q ; a point $G \in E$ of large prime order n (generator of the group of points of order n). Parameters chosen as such are generally believed to offer a security level of $\frac{n}{2}$ given current

knowledge and technologies. Parameters are recommended to be generated following [16]. The field size q is usually taken to be a large, odd prime or a power of 2. The implementation of OpenSSL uses both prime fields and $q = 2^m$ though the results in this paper relate to the binary field case.

Public-Private Key pairs: the private key is an integer d , $1 < d < n - 1$ and the public key is the point $Q = dG$. Calculating the private key from the public key requires solving the ECDLP, which is known to be hard in practice for the correctly chosen parameters. The most efficient algorithms currently known which solve the ECDLP have a square root run time in the size of the group, hence the aforementioned security level.

Suppose Bob, with public-private Key pair $\{d_B, Q_B\}$, wishes to send a signed message m to Alice, he follows the following steps:

1. Using an approved hash algorithm, compute $e = \text{Hash}(m)$, take \bar{e} to be the leftmost n bits of e . ***check fips is leftmost or mod***
2. Randomly select $k \leftarrow_R \mathbb{Z}_n$ with $1 < k < p - 1$ and $(k, p - 1) = 1$.
3. Compute the point $(x, y) = kG \in E$. ***using montgomery ladder or otherwise***
4. Take $r = x \bmod n$; if $r = 0$ then return to step 2.
5. Compute $s = k^{-1}(z + rd_B) \bmod n$; if $s = 0$ then return to step 2.
6. Bob sends (m, r, s) to Alice.

The message m is not necessarily encrypted, the contents may not be secret, but a valid signature gives Alice strong evidence that the message was indeed sent by Bob. She verifies that the message came from Bob by

1. checking that all received parameters are correct, that $r, s \in \mathbb{Z}_n$ and that Bob's public key is valid, that is $Q_b \neq \mathcal{O}$ and $Q_B \in E$ is of order n .
2. Using the same hash function and method as above, compute \bar{e} .
3. Compute $\bar{s} = s^{-1} \bmod n$.
4. Find the point $(x, y) = \bar{e}sG + r\bar{s}Q_B$.
5. Verify that $r = x \bmod n$ otherwise reject the signature.

Step 2 of the signing algorithm is of vital importance, inappropriate reuse of the random integer is what led to the highly publicised breaking of Sony PS3 implementation of ECDSA. Knowledge of the random value k leads to knowledge of the secret key as all values (m, r, s) can be observed by an eavesdropper, \bar{e} can be found from m , $r^{-1} \bmod n$ can be easily found from n , and if k is discovered then an adversary can find Bob's secret key through the simple calculation

$$d_B = (sk - \bar{e})r^{-1}.$$

2.2 The Montgomery Ladder

Scalar multiplication is a common operation in cryptography and in a number of incidences (such as the multiplication by the secret, randomly generated element required in ECDSA), the scalar is intended to remain secret. This scalar

Input: Point P , scalar n , k bits
Output: Point nP
 $Q \leftarrow \mathcal{O}$
for i from k to 0 **do**
 $Q \leftarrow 2Q$ **if** $n_i = 0$ **then**
 $Q \leftarrow Q + P$
 end
end

Algorithm 1: Double-and-Add Point Multiplication

multiplication is most efficiently performed using a square-and-multiply method (or the related Right-to-left method) as outlined in Algorithm 1

Double-and-add methods, though efficient, are vulnerable to simple power analysis. The addition law for points on Weirstrass curves is not complete, that is, the computation of $P + Q$ differs between the cases $P = Q$ and $P \neq Q$. Consequently, by examining the power consumption of the computation it is possible to distinguish when the if loop is executed and hence when a bit of n is 0.

As described by Montgomery in [14]

Input: Point P , scalar n , k bits
Output: Point nP
 $R_0 \leftarrow \mathcal{O}$
 $R_1 \leftarrow P$
for i from k to 0 **do**
 if $n_i = 0$ **then**
 $R_1 \leftarrow R_0 + R_1$
 $R_0 \leftarrow 2R_0$
 else
 $R_0 \leftarrow R_0 + R_1$
 $R_1 \leftarrow 2R_1$
 end
end

Algorithm 2: Montgomery Ladder Point Multiplication

2.3 The Flush+Reload attack

Spatial prefetching limits the spatial resolution of the FLUSH+RELOAD attack. When probing a cache line, FLUSH+RELOAD issues a read from the line. The spatial prefetcher, then, prefetches the pair of the cache line. This prefetching prevents using FLUSH+RELOAD on the pair of the line, forcing the spatial resolution to be lower than a probe per two cache lines.

The wait period between flushing and reloading imposes a limit on the temporal resolution of FLUSH+RELOAD. The order of memory accesses occurring

during a single wait period cannot be observed. Thus, shorter wait periods provide a higher observation resolution. On the other hand, with a short wait, a more

2.4 Related Work

There have been a number of publications addressing the security issues of digital signatures when the nonce is partially leaked, including [6, 7, 17]. These attacks usually rely on having obtained a relatively small number of bits of the ephemeral keys used for many signatures and then use the LLL method [12] to solve the related hidden number to find the secret value. The attack of [17], for example, given a group of order around 160 bits the probabilistic algorithm in would obtain the secret key using 23 signatures (assuming independent and uniformly at random selected messages) in polynomial time, using only 7 consecutive, least significant leaked bits of the nonce (relying on some reasonable assumptions). Each of these assumes only a small fraction of k is recovered. In these works the reconstruction of the secret information from the partial nonce information is the main contribution. The main contribution of this work is to illustrate the method, the adapted technique of [22], to recover a large majority of the bits, from which the full nonce is then obtained using only one signature. Once the nonce has been fully determined the secret key is obtained using only less than one second of additional computation time (idle computer). For example, we are able to obtain around 560 of the 571 bits of a nonce. Though the goal and approach of the works are similar, the methods are very different.

The attack of [3] uses the above methods to highlight specifically a vulnerability in OpenSSL's ladder implementation for curves over binary fields, the same target of our attack. Though the attacks differ, they both illustrate that the OpenSSL implementation of the Montgomery ladder is vulnerable from both remote attacks and attacks launched from virtual machines with access to the memory of the target computer. The countermeasure suggested in [3] will not thwart this attack.

3 Attacking OpenSSL ECDSA

OpenSSL is one of the most common open-source cryptographic libraries. It provides a set of cryptographic services, including public and shared key encryption algorithms and public key signature algorithms.

OpenSSL's implementation of ECDSA uses the Montgomery ladder algorithm for scalar multiplication on the Elliptic Curve. We use this implementation to demonstrate that naïve implementations of the Montgomery ladder are susceptible to the FLUSH+RELOAD attack.

Listing 1.1 shows the relevant section of the implementation of the Montgomery ladder in OpenSSL version 1.0.1e. The bits of the multiplication scalar are stored in the word array `scalar->d`. The outer loop, at lines 268 to 286 traverses over the words representing the scalar. The inner loop, at lines 271

to 284 traverses the bits in each word. Line 273 tests the bit. For each bit the implementation executes a group add followed by a group double. If the bit is set, the implementation uses lines 275 and 276. For clear bits it uses lines 280 and 281.

Listing 1.1. OpenSSL Implementation of the Montgomery Ladder

```

268 for (; i >= 0; i--)
269 {
270     word = scalar->d[i];
271     while (mask)
272     {
273         if (word & mask)
274         {
275             if (!gf2m_Madd(group, &point->X, x1, z1, x2, z2, ctx))
276                 goto err;
277             if (!gf2m_Mdouble(group, x2, z2, ctx)) goto err;
278         }
279         else
280         {
281             if (!gf2m_Madd(group, &point->X, x2, z2, x1, z1, ctx))
282                 goto err;
283             if (!gf2m_Mdouble(group, x1, z1, ctx)) goto err;
284         }
285         mask >>= 1;
286     }
287     mask = BN_TBIT;
288 }

```

As the listing demonstrates, the Implementation is very regular. For each bit, the implementation executes exactly the same sequence of operations. The only differences between set and clear bit are the lines that invoke these operations. While this is a small difference, it is sufficient for mounting an attack that recovers the values of the bits.

Our spy program uses the FLUSH+RELOAD technique to monitor the execution of the `if` statement in line 273. We distinguish between executing the `then` and the `else` blocks of the `if` statement. This information reveals the value of the bit tested by the `if` statement.

FLUSH+RELOAD monitors execution by placing probes on shared memory lines. For FLUSH+RELOAD to recover the bit value, it must distinguish memory lines access sequences that result from a set bit from those resulting from a clear bit. Achieving this depends on several factors: the mapping of source code to memory lines, the sequence of accesses to these memory lines when executing the code and FLUSH+RELOAD's ability to accurately capture the sequences.

The mapping of source lines to cache lines in our build of OpenSSL is depicted in Diagram 1. The machine code created from source lines 273 to 282 covers the virtual memory address range 0x0812130C to 0x081213e8. This range spans four cache lines, marked *A*, *B*, *C* and *D*.

The minimum sequence of memory line accesses required for executing this code can now be constructed. The `if` statement at line 273 is executed for each bit. The code of this statement is in memory line *A*, hence this line is accessed




Fig. 1. Mapping from Source Code to Memory

when processing of a bit starts. For set bit, the processing continues with source line 275, which maps to memory lines *A* and *B*. The actual call to the group add function occurs at address 0x08121347. (See mark in Diagram 1.) After a delay for computing the group add, execution continues in memory line *B* to process the return value and to invoke the group doubling function. The group doubling function returns to memory line *B* and execution leaves the *if* body at memory line *D*.

Hence, the sequence of memory line accesses required for a set bit is: *A*, *B*, *add*, *B*, *double*, *B*, *D*. Similarly, for a clear bit, the sequence is: *A*, *C*, *add*, *C*, *D*, *double*, *D*.

Due to the limited temporal resolution of FLUSH+RELOAD, the attack can observe the order of memor accesses only if they are sufficiently separated in time. Hence, in the case of OpenSSL, the attack can only observe the order of memory accesses if they are separated by a call to a group operation. For example, when the bit is set, the attack cannot decide whether the access to memory line *A* precedes or follows the access to memory line *B*. Similarly, when observed by FLUSH+RELOAD, memory accesses issued after the group double are merged with those issued at the start of processing the following bit. Diagram 2 shows the memory accesses observable by FLUSH+RELOAD when processing a set bit followed by a clear bit.




Fig. 2. Observable Memory Access over Time

The diagram also shows memory accesses issued by processor optimisations. These optimisations pre-load memory lines into the cache to reduce the time the program waits for these lines. For example, when the processor uses speculative execution [20], it follows both arms of a conditional branch before evaluating the condition. When the condition is evaluated, the processor commits to the pre-processed computation of the correct arm, disposing of the computation done for the other arm. In the case of OpenSSL this means that even before evaluating the bit, the processor may start processing both line 275 and line 280, triggering memory loads from memory lines *A*, *B* and *C*.

Another optimisation that can cause additional memory line access is spatial prefetching [8]. The processor pairs adjacent memory lines and tries to bring both memory lines into the cache when there is a miss on one of the pair's line. For example, when there is a cache miss on memory line *A*, the spatial prefetcher may attempt to prefetch memory line *B* and vice versa.

Consequently, as demonstrated in Diagram 2, the memory lines accessed between computing the group add and the group double can be used for recovering the value of the bit. Probing any of lines *A* and *B* gives a positive indication of set bits. Probing any of lines *C* and *D* gives a positive indication of clear bits. For our attack we probe memory lines *B* and *D*. The next section describes the details of our experiment with the attack and its results.

4 Experimental Setup and Results

To test the attack on OpenSSL we used an HP Elite 8300 running Fedora 18. As the openssl shipped with Fedora does not support elliptic curves cryptography, we used our own build of OpenSSL 1.0.1e. To facilitate the mapping from source lines to memory addresses we built OpenSSL with debugging symbols. In a real attack settings, the attacker will need to reverse engineer [5] the OpenSSL library. For the experiment we used the OpenSSL `sect1571r1` curve (NIST Binary-Curve B-571 [16].)

- ★ Operation times
- ★ Choice of slot size
- ★ Example of results
- ★ Number of lost bits

4.1 Full recovery of the nonce

There are a number of attacks on ECC protocols which recover the secret key using a small proportion of the leaked nonce. In this attack, the proportion of the nonce recovered by the FLUSH+RELOAD attack is significantly higher; using one of the existing attacks would be unnecessarily (computationally) excessive. Attacks using lattice techniques to solve the related hidden number problem, for example, ... In experiments a basic BSGS implementation sufficed...

Distribution of the missing nonce bits

5 Discussion

- ★ **LLL attack**
 - ★ **Distribution of missing bits**
analysis for when we are not as successful.
 - ★ **Expected number of observed signatures to break key**

Mitigation

As the ECDLP is not targeted by this attack, the signature protocol it is made no more vulnerable by our results. This attack is targets the scalar multiplication implmentation of OpenSSL and is therefore particular to implementations using

In the Networking and Cryptography library (NaCl), implemented by Daniel J. Bernstein, Tanja Lange and Peter Schwabe, there is no data flow from secrets to branch conditions, precisely the vulnerability of the OpenSSL implementation targeted by this attack. Analysis of the core security features of NaCl is given in [2]. The attack presented in this article relies on distinguishing bits of the nonce by observing the branching in traditional Montgomery ladder implementation. As the NaCl library avoids branching dependent on secret parameters this attack is not applicable to NaCl's `crypto_sign` API. (NaCl is in the public domain and has been made available by the authors of [2] at <http://nacl.cr.yp.to>.)

6 Conclusions and future work

The results of this work and [4] imply that the OpenSSL montgomery ladder implementation should be avoided in all implementations of elliptic curve protocols when a scalar multiplication step involves a secret parameter. This attack is applicable when the malicious party has access to the memory of the targeted device which is not a completely unreasonable assumption ***.

The results of this work also support the theory of [21] that smaller keys are more resilient to side-channel analysis, in this attack a higer proportion of the nonce was obtained for larger key sizes. This implies that as we naturally transition to larger parameters in response to increasing computing capabilities, prevention of side-channel attacks should be incorporated into the implementation design, as is the methodology adopted by the authors of the NaCl cryptographic library.

There are a number of future directions for this research, including the development of the use of Gray codes in improving the BSGS method to extend the efficiency of obtaining the full scalar from a partially obtained secret. It would also be useful to understand what proportion of the nonce should be obtained for the BSGS method to become more practical than the LLL method.

Another related area of interest is how to adapt this attack to recover bits when a sliding window is used; using the method presented in this paper we are able to extract the first and last bit of the sliding window (two bits per window) and thus the proportion of bits obtained decreases as the window size increases.

7 Acknowledgements

The authors wish to thank Dr Katrina Falkner for the advice and support.

This research was performed under contract to the Defence Science and Technology Organisation (DSTO) Maritime Division, Australia.

References

1. American National Standards Institute: ANSI X9.62, Public Key Cryptography for the Financial Services Industry: The Elliptic Curve Digital Signature Algorithm. (1999)
2. Bernstein, D.J., Lange, T., Schwabe, P.: The security impact of a new cryptographic library. In: Proceedings of the 2nd international conference on Cryptology and Information Security in Latin America. LATINCRYPT'12, Berlin, Heidelberg, Springer-Verlag (2012) 159–176
3. Brumley, B.B., Tuveri, N.: Remote timing attacks are still practical. In Atluri, V., Diaz, C., eds.: Computer Security - ESORICS 2011. Volume 6879 of Lecture Notes in Computer Science., Springer-Verlag (2011) 355–371
4. Brumley, B.B., Tuveri, N.: Remote timing attacks are still practical. Cryptology ePrint Archive, Report 2011/232 (2011) <http://eprint.iacr.org/>.
5. Cipresso, T., Stamp, M.: Software reverse engineering. In Stavroulakis, P., Stamp, M., eds.: Handbook of Information and Communication Security. Springer (2010) 659–696
6. Gopalakrishnan, K., Thériault, N., Yao, C.Z.: Solving discrete logarithms from partial knowledge of the key. In Srinathan, K., Pandu Rangan, C., Yung, M., eds.: Progress in Cryptology – INDOCRYPT 2007. Volume 4859 of Lecture Notes in Computer Science., Springer (2007) 224–237
7. Howgrave-Graham, N., Smart, N.P.: Lattice attacks on digital signature schemes. Des. Codes Cryptography **23**(3) (2001) 283–290
8. Intel Corporation: Intel 64 and IA-32 Architecture Optimization Reference Manual. (April 2012)
9. Johnson, D., Menezes, A., Vanstone, S.: The elliptic curve digital signature algorithm (ECDSA). International Journal of Information Security **1**(1) (August 2001) 36–63
10. Joye, M., Yen, S.M.: The Montgomery powering ladder. In Kaliski, Jr., B.S., Koç, Ç.K., Paar, C., eds.: Cryptographic Hardware and Embedded Systems—CHES 2002. Volume 2523 of Lecture Notes in Computer Science., Springer-Verlag (2003) 291–302
11. Koblitz, N.: Elliptic curve cryptosystems. Mathematics of Computation **48**(177) (January 1987) 203–209
12. Lenstra, A., Lenstra, H., Lovász, L.: Factoring polynomials with rational coefficients. Math. Ann. **261** (1982) 515–534
13. Miller, V.S.: Use of elliptic curves in cryptography. In Williams, H.C., ed.: Advances in Cryptology - CRYPTO'85. Volume 218 of Lecture Notes in Computer Science., Springer-Verlag (1986) 417–426
14. Montgomery, P.L.: Speeding the Pollard and elliptic curve methods of factorization. Mathematics of Computation **48**(177) (January 1987) 243–264
15. National Institute of Standards and Technology: Digital Signature Standard, FIPS Publication 186-2. (2000)

16. National Institute of Standards and Technology: FIPS PUB 186-4 Digital Signature Standard (DSS). (2013)
17. Nguyen, P.Q., Shparlinski, I.E.: The insecurity of the elliptic curve digital signature algorithm with partially known nonces. *Designs, Codes and Cryptography* **30**(2) (September 2003) 201–217
18. Okeya, K., Kurumatani, H., Sakurai, K.: Elliptic curves with the Montgomery-form and their cryptographic applications. In: *Public Key Cryptography*. Volume 1751 of *Lecture Notes in Computer Science.*, Springer-Verlag (2000) 238–257
19. OpenSSL. <http://www.openssl.org>
20. Uht, A.K., Sindagi, V.: Disjoint eager execution: An optimal form of speculative execution. In: *Proceedings of the 28th International Symposium on Microarchitecture*, Ann Arbor, Michigan, United States (November 1995) 313–325
21. Walter, C.D.: Longer keys may facilitate side channel attacks. In Matsui, M., Zuccherato, R.J., eds.: *Selected Areas in Cryptography*. Volume 3006 of *Lecture Notes in Computer Science.*, Springer-Verlag (2004) 42–57
22. Yarom, Y., Falkner, K.: Flush+reload: a high resolution, low noise, l3 cache side-channel attack. *Cryptology ePrint Archive*, Report 2013/448 (2013) <http://eprint.iacr.org/>.